

# Introduction to Machine Learning

Alessandro Rudi, Umut Simsekli

Notes by Antoine Groudiev

12th February 2024

## Contents

<b>1</b>	<b>An overview of Machine Learning</b>	<b>1</b>
1.1	What is ML? . . . . .	1
1.2	Topics in Machine Learning . . . . .	1
1.2.1	Supervised Learning . . . . .	1
1.2.2	Probabilistic approach . . . . .	2
1.2.3	Unsupervised Learning . . . . .	2

## Introduction

### 1 An overview of Machine Learning

#### 1.1 What is ML?

Considering a problem, such as image classification: given an input image of a dog or a cat, the program is asked to determine whether the image is a dog or a cat. Conventional programming would hardcode the solution to this problem. But this process takes time and is not easily generalisable. Instead, an ML model is trained on a dataset to produce a program to solve the problem.

Many successful applications of Machine Learning are:

- Face recognition
- Spam filtering
- Speech recognition
- Self-driving systems; pedestrian detection

#### 1.2 Topics in Machine Learning

##### 1.2.1 Supervised Learning

**Example** (Classification). *Features*  $x \in \mathbb{R}^d$ , *labels*  $y \in \{1, \dots, k\}$

**Definition** (Regression). Features  $x \in \mathbb{R}^d$ , labels  $y \in \mathbb{R}$ . To tackle such problem, we look for a parametrized function  $f_\theta(x_i) \simeq y_i$  for some  $f_\theta$  in a function space

$$\mathcal{F} = \{f_\theta : \theta \in \Theta\}$$

Our goal is therefore to find the best function in  $\mathcal{F}$  such that  $f$  "fits" the training data. For example, we can say that  $f$  "fits" the training data when

$$\frac{1}{n} \sum_{i=1}^n (f(x_i) - y_i)^2$$

is "small". Such a function is not interesting in general, like for classification.

**Definition** (Loss function). Assumes that the features are in  $\mathcal{X}$  and the labels are in  $\mathcal{Y}$ . We introduce the more general *loss function* notion:

$$l : \mathcal{Y}^2 \rightarrow \mathbb{R}_+$$

For a regression task, we can use  $l(\hat{y}, y) = (\hat{y} - y)^2$ . For a classification task,  $l(\hat{y}, y) = \mathbb{1}_{\hat{y} \neq y}$ .

Therefore, for a regression problem, we might choose:

$$f^* = \operatorname{argmin}_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n l(f(x_i), y_i)$$

In the parametric case, when  $\mathcal{F} = \{f_\theta : \theta \in \Theta\}$ , we might minimize with respect to  $\theta$ :

$$\theta^* = \operatorname{argmin}_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n l(f_\theta(x_i), y_i)$$

### 1.2.2 Probabilistic approach

Let  $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$  be the feature space. Let  $D$  be a distribution on  $\mathcal{Z}$ ; we make the assumption that the training data is iid from  $D$ :

$$(x_i, y_i) \sim D$$

and the same thing hold for the test data:

$$(\tilde{x}_i, \tilde{y}_i) \sim D$$

According to the Strong Law of Large Numbers, the test loss converges almost surely:

$$\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n l(f_\theta(\tilde{x}_i), \tilde{y}_i) = \mathbb{E}_{(x,y) \sim D} [l(f_\theta(x), y)] =: R(\theta) = R(f_\theta)$$

where  $R(\theta)$  is the *population risk*.

**Definition** (Risk minimization).

### 1.2.3 Unsupervised Learning

**Example** (Clustering).

**Example** (Dimensionality reduction). We are given features  $x \in \mathbb{R}^d$  and labels  $y \in \{0, 1\}$  which form a "training" dataset  $S = \{(x_1, y_1), \dots, (x_n, y_n)\}$ . We assume that  $d \gg 1$ ; our goal is to find  $d' \ll d$  such that  $(x_1, y_1, \dots)$