

Semantics and Applications to Verification

Xavier Rival and Jérôme Feret

Notes by Antoine Groudiev

6th March 2024

Contents

1	Introduction to Semantics	2
1.1	Case studies	2
1.1.1	Ariane 5 – Flight 501	2
1.1.2	Lufthansa Flight 2904	2
1.1.3	Patriot missile (anti-missile system), Dahran	3
1.1.4	General remarks	3
1.2	Approaches to verification	3
1.2.1	The termination problem	3
1.2.2	A summary of common verification techniques	5
1.3	Orderings, lattices and fixpoints	5
1.3.1	Basic definitions on orderings	5
1.3.2	Complete lattice	6
1.3.3	How to prove semantic properties	6
1.3.4	Operators over a poset	6
1.3.5	Fixpoints theorems	7
2	Operational Semantics	7
2.1	Definition and properties	7
2.2	Examples	8
2.2.1	Word recognition	8
2.2.2	Pure λ -calculus	8
2.3	A MIPS like assembly language	9
2.4	Traces semantics	9
3	Traces Properties	9
3.1	A high level overview	9
3.2	Safety properties	10
3.2.1	Informal and formal definitions	10
3.2.2	A few operators on traces	10
3.2.3	Formal definition of safety	11
3.2.4	Intuition of the formal definition	11
3.2.5	Proof method for safety properties	11
3.3	Liveness properties	11
3.3.1	Informal and formal definitions	11
3.3.2	Proof method	11
3.4	Decomposition of trace properties	11
3.5	A specification language: temporal logic	11
3.6	Beyond safety and liveness	11

Introduction

This document is Antoine Groudiev's class notes while following the class *Sémantique et applications à la vérification de programmes* (Semantics and Applications to Verification) at the Computer Science Department of ENS Ulm. It is freely inspired by the class notes of Xavier Rival.

1 Introduction to Semantics

1.1 Case studies

We will study some examples of software errors: what are the causes of these, what kind of properties do we want to verify in order to prevent such failures?

1.1.1 Ariane 5 – Flight 501

Ariane 5 was a satellite launcher, aimed at replacing Ariane 4. Its first flight, June, 4th, 1996, was a failure, with more than \$370 000 000 of damages. 37 seconds after the launch, the rocket exploded.

The system contained sensors, two calculators (SRI, OBC), actuators, and redundant systems (failure tolerant system). The failure was due to an unhandled arithmetic error. Each register of the SRI has a size of 16, 32, or 64 bits. The error was due to a conversion of a 64-bit float to a 16-bit integer. The value was too large to be represented in 16 bits, and the conversion failed. The software was not able to handle this error, and the system crashed.

Several solutions would have prevented this mishapening:

- Desactivate interruptions on overflows
- Fix the SRI code, so that no overflow can happen. All conversions must be *guarded against overflows*:

```
double x = /* ... */ ;
short i = /* ... */ ;
if ( -32768. <= x && x <= 32767. )
    i = ( short ) x ;
else
    i = /* default value */ ;
```

This may be costly, but redundant tests can be removed.

- Handle conversion errors (not trivial): identify the problem and fix it at run time.

The piece of code that generated the error was used to do a useless task, the re-calibration process is not usefull after take-off. Furthermore, the code was already used in Ariane 4; initially protected by a safety guard, many conversions and tests were removed for the sake of performance after being tested on Ariane 4.

The crash was not prevented by redundant systems: the two calculators were running the same code, and the same error was made on both. Redundancy can prevent hardware errors, but is not enough to prevent software errors.

1.1.2 Lufthansa Flight 2904

On November 22, 2003, a Lufthansa Airbus A320-200 crashed at the airport of Warsaw, Poland. The plane was landing, and the weather was bad. The plane was not able to stop before the end of the runway, and crashed into a building. The cause of the crash was a software error

in the plane's computer. The plane was not able to compute the correct deceleration, and the pilot was not able to stop the plane in time.

1.1.3 Patriot missile (anti-missile system), Dahran

The purpose of the Patriot system was to destroy foe missiles before they reach their target, and was used in the first Gulf War with a success rate around 50%. The system was used to destroy Scud missiles, and the system was not able to destroy one of them, which hit a barrack, killing 28 people. The cause of the failure was a software error in the system's clock. The system was not able to compute the time correctly due to fixed precision arithmetic error, and the missile was not destroyed.

1.1.4 General remarks

The examples given so far are not isolated cases. The typical causes of software errors are:

- Improper specification
- Incorrect implementation of a specification (the code should be free of runtime errors, and should produce a result that meets some property)
- Incorrect understanding of the execution model (generation of too imprecise results)

This creates new challenges to ensure embedded systems do not fail. The main techniques to ensure software safety are software development techniques:

- software engineering
- programming rules
- make software cleaner

In this class, we will instead dive into formal methods:

- should have sound mathematical foundations
- should allow guaranteeing software meet some complex properties
- should be trustable
- increasingly used in real life applications

This course will contain two main parts. The first part will be about Semantics, which allow describing precisely the behavior of programs, express the properties to verify, and to transform and compile programs. The second part, Verification, aims at proving semantic properties of programs. A very strong limitation of verification is undecidability; several approaches make various compromises around undecidability.

1.2 Approaches to verification

1.2.1 The termination problem

Definition (Termination). A program P terminates on input X if and only if any execution of P with input X eventually reaches a final state. A final state is a final point in the program (i.e., not an error).

Definition (Termination problem). Can we find a program P_t that takes as argument a program P and data X and that returns **True** if P terminates on X , and **False** otherwise?

Property 1. *The termination problem is not computable.*

Proof. We assume there exists a program Pa such that Pa always terminates, and returns 1 if and only if P terminates on input X . We consider the following program:

```
void P0 ( P ) {
    if ( Pa ( P , P ) == 1 ) {
        while ( 1 ) {
            // loop forever
        }
    } else {
        return ; // do nothing
    }
}
```

and we consider the return value of $Pa(P0, P0)$. If $Pa(P0, P0) == 1$, then $P0$ loops forever, and if $Pa(P0, P0) == 0$, then $P0$ terminates. This is a contradiction, and the termination problem is not computable. \square

Property 2. *The absence of runtime errors is not computable. We cannot find a program Pc that takes a program P and input X as arguments, always terminates, and returns 1 if and only if P runs on input X without a runtime error.*

Theorem (Rice theorem). *Considering a Turing complete language, any non-trivial semantic specification is not computable. Therefore, all interesting properties are not computable (termination, absence of runtime/arithmetic errors, etc.).*

The initial verification problem is not computable. Several compromises can be made: simulation, testing, assisted theorem proving, model checking, bug-finding, static analysis with abstraction.

Definition (Safety verification problem). The Semantics $\llbracket P \rrbracket$ of a program P is the set of behaviors of P (e.g. states). A property to verify \mathcal{S} is the set of admissible behaviors (e.g. safe states). Our goal is to establish $\llbracket P \rrbracket \subseteq \mathcal{S}$

$\llbracket P \rrbracket$ can be sound (identify any wrong program), complete (accept all correct programs), and automated, but not all three at the same time.

Testing by simulation The principle of testing by simulation is to run the program on finitely many finite inputs, to maximize coverage and inspect erroneous traces to fix bugs. It is very widely used, through unit testing, integration testing, etc. It is both automated and complete, but is unsound and costly.

Machine assisted proof The principle of machine assisted proof is to have a machine check proof that is partly human written: tactics or solvers may help in the inference, and the hardest invariants have to be user-supplied. It is sound and quasi-complete, but not fully automated and costly.

Model checking We consider finite systems only, using algorithms for exhaustive exploration, symmetry reduction, ... It is automated, sound, and complete *with respect to the model*.

Bug finding The principle of bug finding is to identify "likely" issues, patterns known to often indicate an error: it uses bounded symbolic execution, model exploration, and rank "defect" reports using heuristics. It is neither sound nor complete, but is fully automated.

Static analysis with abstraction The principle of static analysis with abstraction is to use some approximation, but always in a conservative manner. We can use under-approximation of the property to verify:

$$\mathcal{S}_{\text{under}} \subseteq \mathcal{S}$$

and over-approximation of the semantics:

$$\llbracket P \rrbracket \subseteq \llbracket P \rrbracket_{\text{upper}}$$

We let an automatic static analyser attempt to prove that:

$$\llbracket P \rrbracket_{\text{upper}} \subseteq \mathcal{S}_{\text{under}}$$

If it succeeds, then we have proven that $\llbracket P \rrbracket \subseteq \mathcal{S}$. It is automated, sound, but incomplete.

1.2.2 A summary of common verification techniques

	Automatic	Sound	Complete	Source level
Simulation	Yes	No	Yes	Yes
Assisted Proving	No	Yes	Almost	Partially
Model-checking	Yes	Yes	Partially	No
Bug-finding	Yes	No	No	Yes
Static analysis	Yes	Yes	No	Yes

1.3 Orderings, lattices and fixpoints

1.3.1 Basic definitions on orderings

Definition (Partially ordered set (poset)). Let a set \mathcal{S} and a binary relation $(\sqsubseteq) \subseteq \mathcal{S} \times \mathcal{S}$ over \mathcal{S} . Then, \sqsubseteq is an order relation if and only if it is reflexive, transitive, antisymmetric. Furthermore, we define $x \sqsubset y ::= (x \sqsubseteq y \wedge x \neq y)$. Most orders in this class won't be total.

We often use Hasse diagrams to represent posets.

In the following, we illustrate order relations and their usefulness in semantics using the standard definition of *word automata*. The semantics of an automaton is the set of words recognized by it.

We can already define a few semantic properties:

- \mathcal{P}_0 : no recognized word contains two consecutive b

$$\mathcal{L}[\mathcal{A}] \subseteq L^* \setminus L^*bbL^*$$

- \mathcal{P}_1 : all recognized words contain at least one occurrence of a

$$\mathcal{L}[\mathcal{A}] \subseteq L^*aL^*$$

- \mathcal{P}_2 : recognized words do not contain b

$$\mathcal{L}[\mathcal{A}] \subseteq (L \setminus \{b\})^*$$

Definition (\perp , \top). When they exist, we denote by infimum \perp and supremum \top the smallest and largest elements of the poset.

Definition (\sqcap , \sqcup). We denote $\sqcap \mathcal{S}$ the glb (greatest lower bound) of \mathcal{S} , and \sqcup the lub (lowest upper bound) of \mathcal{S} .

1.3.2 Complete lattice

Definition (Complete lattice). A complete lattice is a tuple $(\mathcal{S}, \sqsubseteq, \perp, \top, \sqcup, \sqcap)$ where $(\mathcal{S}, \sqsubseteq)$ is a poset, and any subset \mathcal{S}' of \mathcal{S} has a glb $\sqcap \mathcal{S}'$ and a lub $\sqcup \mathcal{S}'$.

Definition (Lattice). A lattice is a tuple $(\mathcal{S}, \sqsubseteq, \perp, \top, \sqcup, \sqcap)$ where $(\mathcal{S}, \sqsubseteq)$, and any pair $\{x, y\}$ of \mathcal{S} has a glb $x \sqcap y$ and a lub $x \sqcup y$.

Example. $\mathbb{Q} \cap [0, 1]$ is a lattice but not a complete lattice, since

$$\left\{ q \in \mathbb{Q} \cap [0, 1] \mid q \leq \frac{\sqrt{2}}{2} \right\} \subseteq \mathbb{Q} \cap [0, 1]$$

has no lowest upper bound in $\mathbb{Q} \cap [0, 1]$.

Property 3. A finite lattice is also a complete lattice.

Definition (Increasing chain). Let $(\mathcal{S}, \sqsubseteq)$ be a poset and $\mathcal{C} \subseteq \mathcal{S}$. It is an increasing chain if and only if it is not empty, and $(\mathcal{C}, \sqsubseteq)$ is total.

Example. In the powerset $(\mathcal{P}(\mathbb{N}), \subseteq)$,

$$\mathcal{C} := \left\{ \{2^0, \dots, 2^i\} \mid i \in \mathbb{N} \right\}$$

is an increasing chain.

Definition (Increasing chain condition). The poset $(\mathcal{S}, \sqsubseteq)$ satisfies the increasing chain condition if and only if any increasing chain $\mathcal{C} \subseteq \mathcal{S}$ is finite.

Definition (Complete partial order). A complete partial order (cpo) is a poset $(\mathcal{S}, \sqsubseteq)$ such that any increasing chain \mathcal{C} of \mathcal{S} has at least an upper bound. A pointed cpo is a cpo with a bottom element \perp .

1.3.3 How to prove semantic properties

1.3.4 Operators over a poset

Definition (Operators and orderings). Let $(\mathcal{S}, \sqsubseteq)$ be a poset and $\varphi : \mathcal{S} \rightarrow \mathcal{S}$ be an operator over \mathcal{S} . Then, φ is:

- *monotone* if and only if $x \sqsubseteq y \Rightarrow \varphi(x) \sqsubseteq \varphi(y)$
- *continuous* if and only if, for any **chain** $\mathcal{S}' \subseteq \mathcal{S}$, then:

$$\begin{cases} \text{if } \sqcup \mathcal{S}' \text{ exists, so does } \sqcup \varphi(\mathcal{S}') \\ \text{and } \sqcup \varphi(\mathcal{S}') = \varphi(\sqcup \mathcal{S}') \end{cases}$$

- \sqcup -preserving if and only if:

$$\forall \mathcal{S}' \subseteq \mathcal{S}, \begin{cases} \text{if } \sqcup \mathcal{S}' \text{ exists, so does } \sqcup \varphi(\mathcal{S}') \\ \text{and } \sqcup \varphi(\mathcal{S}') = \varphi(\sqcup \mathcal{S}') \end{cases}$$

Property 4 (Continuity \Rightarrow monotonicity). If φ is continuous, then it is monotone.

Property 5 (\sqcup -preserving \Rightarrow monotonicity). If φ preserves \sqcup , then it is monotone.

1.3.5 Fixpoints theorems

Definition (Fixpoints). Let $(\mathcal{S}, \sqsubseteq)$ be a poset and $\varphi : \mathcal{S} \rightarrow \mathcal{S}$ be an operator over \mathcal{S} .

- A fixpoint of φ is an element x such that $\varphi(x) = x$.
- A pre-fixpoint of φ is an element x such that $x \sqsubseteq \varphi(x)$.
- A post-fixpoint of φ is an element x such that $\varphi(x) \sqsubseteq x$.
- The least fixpoint of φ is a fixpoint x such that, for any fixpoint y , $x \sqsubseteq y$.
- The greatest fixpoint of φ is a fixpoint x such that, for any fixpoint y , $y \sqsubseteq x$.

Theorem (Tarski's). Let $(\mathcal{S}, \sqsubseteq, \perp, \top, \sqcup, \sqcap)$ be a complete lattice and $\varphi : \mathcal{S} \rightarrow \mathcal{S}$ be a monotone operator. Then:

- φ has a least fixpoint $\text{lfp } \varphi$ and $\text{lfp } \varphi = \sqcap \{ x \in \mathcal{S} \mid \varphi(x) \sqsubseteq x \}$
- φ has a greatest fixpoint $\text{gfp } \varphi$ and $\text{gfp } \varphi = \sqcup \{ x \in \mathcal{S} \mid x \sqsubseteq \varphi(x) \}$
- the set of fixpoints of φ is a complete lattice

Example. We consider a set \mathcal{E} , and a subset $\mathcal{A} \subseteq \mathcal{E}$. We let:

$$\begin{aligned} f : \mathcal{P}(\mathcal{E}) &\rightarrow \mathcal{P}(\mathcal{E}) \\ X &\mapsto X \cup \mathcal{A} \end{aligned}$$

According to Tarski's theorem, the smallest fixpoint of f is \mathcal{A} , and the greatest is \mathcal{E} .

Theorem (Kleene's). Let $(\mathcal{S}, \sqsubseteq, \perp)$ be a pointed cpo and $\varphi : \mathcal{S} \rightarrow \mathcal{S}$ be a continuous operator over \mathcal{S} . The φ has a least fixpoint, and

$$\text{lfp } \varphi = \bigsqcup_{n \in \mathbb{N}} \varphi^n(\perp)$$

2 Operational Semantics

Operational semantics are mathematical descriptions of the executions of a program. It is based on a model of programs, called transition systems.

2.1 Definition and properties

Definition (Transition systems (TS)). A transition system is a tuple $(\mathbb{S}, \rightarrow)$ where \mathbb{S} is the set of states of the system, and $\rightarrow \subseteq \mathbb{S} \times \mathbb{S}$ is the transition relation of the system.

Note that the set of states may be infinite. The majority of interesting examples come from the cases where \mathbb{S} is infinite.

Definition (Deterministic system). A deterministic system is such that a state fully determines the next state.

$$\forall s_0, s_1, s'_1 \in \mathbb{S}, (s_0 \rightarrow s_1 \wedge s_0 \rightarrow s'_1) \implies s_1 = s'_1$$

Otherwise, a transition system is non-deterministic.

The transition relation \rightarrow defines atomic execution steps. It is often called *small-step semantics* or *structured operational semantics*. Steps are *discrete*, and we do not consider non-deterministic systems with probability on transitions (probabilistic transition systems).

Definition (Initial and final states). We often consider transition systems with a set of initial and final states:

1. a set of initial states $\mathbb{S}_I \subseteq \mathbb{S}$ denotes states where the execution should start
2. a set of final states $\mathbb{S}_F \subseteq \mathbb{S}$ denotes states where the execution should reach the end of the program

When needed, we add these to the definition of the transition systems $(S, \rightarrow, \mathbb{S}_I, \mathbb{S}_F)$.

Definition (Blocking state \neq final state). A state $s_0 \in \mathbb{S}$ is blocking when it is the origin of no transition:

$$\forall s_1 \in \mathbb{S}, \neg(s_0 \rightarrow s_1)$$

As an example, we often introduce an error state (usually noted Ω)

2.2 Examples

2.2.1 Word recognition

We can formalize the *word recognition* by a finite automaton using a transition system. We consider an automaton $\mathcal{A} = (Q, q_i, q_f, \rightarrow)$. A state is defined by the remaining of the word to recognize, and the automaton state that has been reached so far. Thus,

$$\mathbb{S} = Q \times \Sigma^*$$

We define the transition relation \rightarrow to be:

$$(q_0, aw) \rightarrow (q_1, w) \iff q_0 \xrightarrow{a} q_1$$

The initial and final states are defined by:

$$\begin{cases} \mathbb{S}_I = \{ (q_i, w) \mid w \in \Sigma^* \} \\ \mathbb{S}_F = \{ (q_f, \varepsilon) \} \end{cases}$$

2.2.2 Pure λ -calculus

A bare-bones model of functional programming:

Definition (λ -terms). The set of λ -terms is defined by:

$$\begin{array}{ll} t, u, \dots ::= x & \text{variable} \\ \mid \lambda x \cdot t & \text{abstraction} \\ \mid tu & \text{application} \end{array}$$

Definition (β -reduction).

The λ -calculus defines a transition system. \mathbb{S} is the set of λ -terms and \rightarrow_β the transition relation. \rightarrow_β is non-deterministic, since multiple β -reduction are sometimes possible. Given a lambda term t_0 , we may consider $(\mathbb{S}, \rightarrow_\beta, \mathbb{S}_I)$ where $\mathbb{S}_I = \{t_0\}$. Blocking states are terms with no redex $(\lambda x \cdot u)v$.

2.3 A MIPS like assembly language

We now consider a very simplified assembly language, containing machine integers \mathbb{B}^{32} . Instructions are encoded over 32 bits and stored in the same space as data, \mathbb{B}^{32} . We assume a fixed set of addresses A .

The memory configuration contains the program counter pc , the general purpose register r_0, \dots, r_{31} , and the main memory (RAM):

$$\text{mem} : A \subseteq \mathbb{B}^{32} \rightarrow \mathbb{B}^{32}$$

Definition (State). A state is a tuple (π, ρ, μ) which comprises a program counter value $\pi \in \mathbb{B}^{32}$, a function mapping each general purpose register to its value $\rho : \{0, \dots, 32\} \rightarrow \mathbb{B}^{32}$, and a function mapping each memory cell to its value, $\mu : A \rightarrow \mathbb{B}^{32}$.

We can define transition relations.

We now look at a more classical imperative language (a bare-bone subset of C), containing variables X , labels L , and values V . A syntax can be defined.

2.4 Traces semantics

Definition (Traces). A *finite trace* is a finite sequence of states s_0, \dots, s_n noted $\langle s_0, \dots, s_n \rangle$. An infinite trace is an infinite sequence of states $\langle s_0, \dots \rangle$. Besides, we write \mathbb{S}^* for the set of finite traces, \mathbb{S}^ω for the set of infinite traces, and $\mathbb{S}^\alpha = \mathbb{S}^* \cup \mathbb{S}^\omega$.

Definition (Concatenation operator \cdot).

3 Traces Properties

3.1 A high level overview

The goal of verification is to prove that $\llbracket P \rrbracket \subseteq \mathcal{S}$ (i.e. that all behaviors of P satisfy specifications \mathcal{S}) where $\llbracket P \rrbracket$ is the program semantics and \mathcal{S} the desired specification. Today, we will mostly focus on program's properties, \mathcal{S} . We will see different families of properties, proof techniques, and specification of properties (are there languages to describe properties?).

A property is a set of traces, defining the admissible executions. There are *safety properties*: something will never happen, which is often proven by invariance; *liveness properties*: something will eventually happen, proven by variance; and beyond safety and liveness, there are *hyperproperties*.

As usual, we consider $\mathcal{S} = (\mathbb{S}, \rightarrow, \mathbb{S}_I)$.

Definition (Properties as sets of states). A property $\mathcal{P} \subseteq \mathbb{S}$. \mathcal{P} if all reachable states belong to \mathcal{P} .

This is the case of the absence of runtime errors, and non-termination.

Definition (Properties as sets of traces). A property \mathcal{T} is a set of traces $\mathcal{T} \subseteq \mathbb{S}^\alpha$. \mathcal{T} if and only if all traces belong to \mathcal{T} , i.e. $\llbracket \mathbb{S} \rrbracket^\alpha \subseteq \mathcal{T}$

State properties are trace properties. Functional properties and termination are trace properties.

Property 6 (Monotonicity).

3.2 Safety properties

3.2.1 Informal and formal definitions

Definition (Informal definition of safety properties). A safety property is a property which specifies that some (bad) behavior defined by a finite, irrecoverable observation will never occur, at any time.

Example. *The following properties are safety properties:*

- Absence of runtime errors
- State properties (the “bad thing” is reaching $\mathbb{S} \setminus \mathcal{T}$)
- Non-termination
- “Not reaching state b after visiting state a ”

Termination is **not** a safety property, since no finite execution is a counter-example of its termination.

We now intend to provide a formal definition of safety. How to refute a safety property? We assume \mathcal{S} does not satisfy safety property \mathcal{P} . Thus, there exists a counter-example trace $\sigma = \langle s_0, \dots, s_n, \dots \rangle \in \llbracket \mathcal{S} \rrbracket \setminus \mathcal{P}$. At this point of our study, the trace may be finite or infinite. The intuitive definition says this trace *eventually exhibits some bad behavior*, that is *irrecoverable* at some *observed at some given time*, thus the observation corresponds to some index i . Therefore, trace $\sigma' = \langle s_0, \dots, s_i \rangle$ violates \mathcal{P} , i.e. $\sigma' \notin \mathcal{P}$. Due to the irrecoverability of the observation, the same goes for any trace with the same prefix. We remark that σ' is finite.

A safety property that does not hold can always be refuted with a finite, irrecoverable counter-example.

3.2.2 A few operators on traces

Definition (Prefix). We write $\sigma_{\upharpoonright i}$ for the prefix of length i of trace σ .

Definition (Suffix or tail). We write $\sigma_{i\downarrow}$ for the suffix of length i of trace σ .

Definition (Upper closure operators (PCI)). In a preorder $(\mathcal{S}, \sqsubseteq)$, a function $\varphi : \mathcal{S} \rightarrow \mathcal{S}$ is an upper closure operator if and only if it is monotone, extensive ($\forall x \in \mathcal{S}, x \sqsubseteq \varphi(x)$) and idempotent.

Definition (Prefix closure). The prefix closure operator is defined by:

$$\begin{aligned} \text{PCI} : \mathcal{P}(\mathbb{S}^\alpha) &\rightarrow \mathcal{P}(\mathbb{S}^\star) \\ X &\mapsto \left\{ \sigma_{\upharpoonright i} \mid \sigma \in X, i \in \mathbb{N} \right\} \end{aligned}$$

PCI is monotone, idempotent, but not extensive on $\mathcal{P}(\mathbb{S}^\alpha)$ (infinite traces do not appear anymore). Its restriction to $\mathcal{P}(\mathbb{S}^\star)$ is extensive.

Definition (The Lim operator). The limit operator is defined by:

$$\begin{aligned} \text{Lim} : \mathcal{P}(\mathbb{S}^\alpha) &\rightarrow \mathcal{P}(\mathbb{S}^\alpha) \\ X &\mapsto X \cup \left\{ \sigma \in \mathbb{S}^\alpha \mid \forall i \in \mathbb{N}, \sigma_{\upharpoonright i} \in X \right\} \end{aligned}$$

Note that the operator Lim is an upper-closure operator.

Proof. Left as an exercise! □

Example. Assume that:

$$\mathcal{S} = \{\varepsilon, \langle a \rangle, \langle a, b \rangle, \langle a, b, a \rangle, \langle a, b, a, b \rangle, \langle a, b, a, b, a \rangle, \dots\}$$

then,

$$\text{Lim}(\mathcal{S}) = \mathcal{S} \uplus \{\langle a, b, a, b, a, b, \dots \rangle\}$$

3.2.3 Formal definition of safety

Definition (The Safe operator). Operator Safe is defined by

$$\text{Safe} = \text{Lim} \circ \text{PCl}$$

Note that Safe is an upper closure operator over $\mathcal{P}(\mathbb{S}^\alpha)$.

Definition (Safety property). A trace property \mathcal{T} is a safety property if and only if it is a fixpoint of the Safe operator, that is

$$\text{Safe}(\mathcal{T}) = \mathcal{T}$$

Furthermore, if \mathcal{T} is a trace property, then $\text{Safe}(\mathcal{T})$ is a safety property, since Safe is idempotent.

Theorem. *Any state property is also a safety property.*

Proof. TBD. □

3.2.4 Intuition of the formal definition

Operator Safe saturates a set of traces S with prefixes and infinite traces of all finite prefixes of which can be observed in S . Thus, if $\text{Safe}(S) = S$ and σ is a trace, to establish that $\sigma \notin S$, it is sufficient to discover a *finite prefix* of σ that cannot be observed in S .

Alternatively, if all finite prefixes of σ belong to S or can be observed as a prefix of another trace in S , by definition of the limit operator, σ belongs to S , *even if it is infinite*. Therefore, the definition captures properties that *can be disproved with a finite counter-example*.

3.2.5 Proof method for safety properties

We consider transition system $\mathcal{S} = (\mathbb{S}, \rightarrow, \mathbb{S}_\mathcal{I})$ and a safety property \mathcal{T} . Finite traces semantics is the least fixpoint of F_* . We seek a way of verifying that \mathcal{S} satisfies \mathcal{T} , i.e. that $\llbracket \mathcal{S} \rrbracket^\alpha \subseteq \mathcal{T}$.

Definition (Invariance proofs). Let \mathbb{I} be a set of finite traces; it is said to be an *invariant* if and only if:

$$\begin{cases} \forall s \in \mathbb{S}_\mathcal{I}, \langle s \rangle \in \mathbb{I} \\ F_*(\mathbb{I}) \subseteq \mathbb{I} \end{cases}$$

3.3 Liveness properties

3.3.1 Informal and formal definitions

3.3.2 Proof method

3.4 Decomposition of trace properties

Theorem. *Let $\mathcal{T} \subseteq \mathbb{S}^\alpha$; it can be decomposed into the conjunction of safety property $\text{Safe}(\mathcal{T})$ and liveness property $\text{Live}(\mathcal{T})$:*

$$\mathcal{T} = \text{Safe}(\mathcal{T}) \cap \text{Live}(\mathcal{T})$$

3.5 A specification language: temporal logic

3.6 Beyond safety and liveness